

Docker Compose Watch for Development in Danswer

This documentation details the steps required to configure a development environment with Docker Compose watch mode in the Danswer application. It includes setting up `docker-compose.dev.yml` for development, creating a development-specific Dockerfile, and with given options like `rebuild`, `sync` and `sync+restart`. Each step ensures efficient live syncing of code changes and an optimal development workflow.

Server Implemented on:

(<http://ec2-54-79-231-211.ap-southeast-2.compute.amazonaws.com>)

(Site: <http://54.79.231.211/>)

(Dev Site: <http://localhost:3000>)

1) `docker-compose.dev.yml`:

Path: `cd danswer/deployment/docker_compose/docker-compose.dev.yml`

Summary: This file defines the services and configuration specifically for the development environment. The setup allows live syncing of changes without the need to rebuild the entire image each time. The watch feature was implemented for the `web_server` service to ensure that all frontend changes are instantly reflected on the live site. I experimented with different watch actions, including `sync`, `sync+restart`, and `rebuild`, to find the most optimal setup. After testing, I determined that the `sync` method is the best choice for frontend development, as it efficiently updates changes without interrupting the service.

Code with action set to `sync`:

`web_server:`

`image: danswer/danswer-web-server:${IMAGE_TAG:-latest}`

`build:`

`context: ../../web`

`dockerfile: Dockerfile.dev - - > (Currently Targeting to Dockerfile.dev within the web dir not to Dockerfile)`

`args:`

`-`

`NEXT_PUBLIC_DISABLE_STREAMING=${NEXT_PUBLIC_DISABLE_STREAMING:-false}`

`-`

`NEXT_PUBLIC_NEW_CHAT_DIRECTS_TO_SAME_PERSONA=${NEXT_PUBLIC_NEW_CHAT_DIRECTS_TO_SAME_PERSONA:-false}`

```

-
NEXT_PUBLIC_POSITIVE_PREDEFINED_FEEDBACK_OPTIONS=${NEXT_PUBLIC
_POSITIVE_PREDEFINED_FEEDBACK_OPTIONS:-}
-
NEXT_PUBLIC_NEGATIVE_PREDEFINED_FEEDBACK_OPTIONS=${NEXT_PUBLIC
_NEGATIVE_PREDEFINED_FEEDBACK_OPTIONS:-}
- NEXT_PUBLIC_DISABLE_LOGOUT=${NEXT_PUBLIC_DISABLE_LOGOUT:-}
-
NEXT_PUBLIC_DEFAULT_SIDEBAR_OPEN=${NEXT_PUBLIC_DEFAULT_SIDEBAR
_OPEN:-}

# Enterprise Edition only
- NEXT_PUBLIC_THEME=${NEXT_PUBLIC_THEME:-}
# DO NOT TURN ON unless you have EXPLICIT PERMISSION from Danswer.
-
NEXT_PUBLIC_DO_NOT_USE_TOGGLE_OFF_DANSWER_POWERED=${NEXT_PU
BLIC_DO_NOT_USE_TOGGLE_OFF_DANSWER_POWERED:-false}
depends_on:
- api_server
restart: always
environment:
- INTERNAL_URL=http://api_server:8080
- WEB_DOMAIN=${WEB_DOMAIN:-}
- THEME_IS_DARK=${THEME_IS_DARK:-}
- DISABLE_LLM_DOC_RELEVANCE=${DISABLE_LLM_DOC_RELEVANCE:-}

# Enterprise Edition only
-
ENABLE_PAID_ENTERPRISE_EDITION_FEATURES=${ENABLE_PAID_ENTERPRIS
E_EDITION_FEATURES:-true}
-
NEXT_PUBLIC_CUSTOM_REFRESH_URL=${NEXT_PUBLIC_CUSTOM_REFRESH_
URL:-}
develop:
watch:
- action: sync
path: ../../web
target: /app
ignore:
- node_modules/
- action: rebuild
path: ../../web/package.json

```

Code with action set to rebuild:

```
develop:
  watch:
    - action: rebuild
      path: ../../web
    ignore:
      - node_modules/
```

Code with action set to sync+restart:

```
develop:
  watch:
    - action: sync+restart
      path: ../../web
      target: /app
    ignore:
      - node_modules/
    - action: rebuild
      path: ../../web/package.json
```

- **Rebuild:** While rebuild updates the `web_server` container, it requires Docker to recreate the container each time whenever a modification is made, which can take a bit of time. This delay makes it less ideal for fast-paced development workflows, as it disrupts the flow by causing a noticeable refresh time.
- **Sync+Restart:** This method, which syncs changes and restarts the container, is only necessary when making changes to both frontend and backend files that are interdependent. For example, if the backend is set up with a multi-stage build process where files from the frontend (such as assets) are copied into the backend at the build stage (using a builder stage), then sync+restart can ensure that backend dependencies on frontend code are accurately reflected. However, for frontend-only updates, sync+restart is an unnecessary step that introduces avoidable delay.
- In summary, sync alone is the optimal method for frontend development in Docker, providing real-time updates without the added delay of container restarts or rebuilds.

2) Dockerfile.dev:

Path: cd danswer/web/Dockerfile.dev

This Dockerfile is tailored specifically for development. It omits production-related optimizations from the original Dockerfile to enable quicker builds and live updates.

Code:

Development Dockerfile (Dockerfile.dev)

Use the base node image for development

FROM node:20-alpine AS base

LABEL com.danswer.maintainer="founders@danswer.ai"

**LABEL com.danswer.description="This image is the web/frontend container of Danswer which **

**contains code for both the Community and Enterprise editions of Danswer. If you do not **

**have a contract or agreement with DanswerAI, you are not permitted to use the Enterprise **

**Edition features outside of personal development or testing purposes. Please reach out to **

founders@danswer.ai for more information. Please visit <https://github.com/danswer-ai/danswer>"

Set up the working directory

WORKDIR /app

Copy only package.json and package-lock.json initially to leverage Docker cache

COPY package.json package-lock.json ./

Install dependencies

RUN npm install

Copy the entire source code

COPY . .

Disable Next.js telemetry for development

ENV NEXT_TELEMETRY_DISABLED=1

Set environment variables for development mode, adjusting as needed

```
ARG NEXT_PUBLIC_DISABLE_STREAMING
ENV
NEXT_PUBLIC_DISABLE_STREAMING=${NEXT_PUBLIC_DISABLE_STREAMING}
```

```
ARG NEXT_PUBLIC_NEW_CHAT_DIRECTS_TO_SAME_PERSONA
ENV
NEXT_PUBLIC_NEW_CHAT_DIRECTS_TO_SAME_PERSONA=${NEXT_PUBLIC_NEW_CHAT_DIRECTS_TO_SAME_PERSONA}
```

```
# Run the development server
CMD ["npm", "run", "dev"]
```

3) *Start and Build Services:*

To start your services with the new development configuration, use the following command. This command builds and recreates the containers from the docker-compose.dev.yml configuration.

Code:

```
sudo docker compose -f docker-compose.dev.yml -p danswer-stack up -d --build
--force-recreate
```

4) *Enable Docker Watch Mode for Live Syncing:*

Docker's watch feature enables live syncing of changes, making it easier to see updates on the site without needing to rebuild the image.

Code:

```
sudo docker compose -f docker-compose.dev.yml -p danswer-stack watch
```

5) *Stop and Remove All Docker Containers:*

To clean up all running containers and remove any stopped containers, use these commands. This is helpful when you need to free up memory or reset the environment.

Code:

- `sudo docker stop $(sudo docker ps -aq)`
- `sudo docker rm $(sudo docker ps -aq)`

6) **Troubleshooting: Memory Shortage:**

- During the setup process, I encountered memory shortages due to high memory consumption by Hot Reloading via Watch feature. Here are a few tips to manage memory usage effectively:
- **Limit Memory Usage in Docker Compose:** You can set memory limits for each container in `docker-compose.dev.yml` to prevent Docker from consuming too much memory.
- Another useful tip when configuring the `sync` method in `docker-compose.dev.yml` is to specify only the files or directories you need to sync, rather than the entire project directory. This can significantly improve performance by reducing the load on Docker's file monitoring system.

7) **Compose.yaml:**

Path: `cd danswer/deployment/docker_compose/compose.yaml`

In our development setup for Danswer, I began with the original `compose.yaml` file provided in the Danswer GitHub repository. This file defines the primary services required for running the application, including `api_server`, `web_server`, `background`, `relational_db`, `index`, and others necessary for the backend infrastructure. The initial `compose.yaml` configuration was meant for general development and testing of core services.

8) **Raised Queries on Docker Community & Danswer Community Github:**

- <https://forums.docker.com/t/dockerized-web-server-not-reflecting-code-changes-despite-sync-confirmation-using-docker-compose-watch/144832>
- <https://github.com/danswer-ai/danswer/issues/3054>

To access the Danswer application as an admin, please log in as:

Email: noelshallum@gmail.com

Password: 12345

- Ibrahim Sultan